

SSLM: Achieving Useful Agent Capabilities at GPT-2 Scale

A Super Small But Usable Architecture

Zice Wang
wangzcl@mails.neu.edu.cn
NTAIRS, Northeastern University, China

May 21, 2026

Abstract

We present SSLM, a hybrid decoder-only language model designed to achieve **useful agent-level capabilities at GPT-2 scale (1.5B parameters)**. The architecture combines three proven innovations: Multi-head Latent Attention (MLA) from DeepSeek V2/V3 for compressed KV-cache and multi-turn conversation memory efficiency, Mamba2 state-space layers for efficient long-context processing and agent planning, and an aux-loss-free Mixture-of-Experts (MoE) module for parameter-efficient scaling without load-balancing auxiliary losses. SSLM v2 achieves **1.45B total parameters** (comparable to GPT-2 XL) with only $\sim 768\text{M}$ active parameters per forward pass, making it well-suited for efficient agent inference. We present complete training results from 22K vocab (SSLM v1) to 53K vocab (SSLM v2), alongside evaluations on C-Eval, TruthfulQA, and GSM8K benchmarks.

1 Introduction

1.1 Goal: Useful Agent at GPT-2 Scale

The original GPT-2 models [Radford et al., 2019] (117M / 345M / 774M / 1.5B) established that decoder-only transformers at the $\sim 1\text{B}$ parameter scale can exhibit surprising language understanding and generation capabilities. Our goal is to match or exceed GPT-2’s utility while being more parameter-efficient, by combining modern architectural innovations (MLA, Mamba2, MoE) that were not available in 2019.

“Useful agent level” means the model should be capable of:

1. **Tool use:** Calling external tools/APIs in a coherent sequence (ReAct-style) [Yao et al., 2023]
2. **Multi-step reasoning:** Breaking down complex tasks into subgoals and executing them
3. **Long-horizon memory:** Maintaining context across extended conversations (≥ 1024 tokens)
4. **Multilingual tool navigation:** Operating on both English and Chinese inputs
5. **Code generation & execution:** Producing correct code for simple tasks

This goal differentiates DenseSLM4MoE from pure benchmark-optimized models — we prioritize architectural choices that enable agent workflows, not just high C-Eval scores.

1.2 Why Hybrid MLA + Mamba2 + MoE?

For agent workloads, three capabilities matter:

Table 1: Architectural solutions for agent capabilities

Capability	Requirement	Solution
Long context memory	KV-cache grows with turns	MLA: compresses KV to 512-dim latent
Efficient processing	48 layers needed	Mamba2: O(N) SSM for 75% of layers
Parameter scale	1B+ params, manageable compute	MoE: 64 experts, top-2 active

1.3 Contributions

1. **Agent-oriented hybrid architecture:** 3:1 Mamba2:MLA layer pattern — Mamba2 layers handle efficient sequence processing; MLA layers provide high-quality attention with compressed KV-cache for multi-turn memory.
2. **Aux-loss-free MoE routing:** No load-balancing auxiliary loss; instead, a learnable `e_score_correction_bias` updated via simple load-based feedback, inspired by DeepSeek V3 [DeepSeek-AI, 2024] and Nemotron-H [NVIDIA, 2024].
3. **Grouped top-k routing:** $n_group = 8$, $topk_group = 2$ for coarse-to-fine expert selection, reducing routing decision overhead.
4. **Complete agent-focused training pipeline:** Muon optimizer [Jordan et al., 2024], WCL learning rate schedule, multilingual pretraining on English/Chinese/code data, SwanLab tracking.
5. **Ablation study** showing that wider experts (8 experts, top-2) outperform narrower experts (32 experts, top-8) at equal parameter count for our target tasks.

2 Architecture

2.1 Design Rationale for Agent Use Cases

For agent workloads, we made the following architectural decisions:

Why 3:1 Mamba2:MLA?

- Mamba2 handles the bulk of forward pass computation (75% of layers) at O(N) complexity, making long-context generation efficient.
- MLA layers appear every 4th layer to provide full attention capability where needed for agentic decisions (e.g., tool selection, instruction following).
- MoE only replaces FFN in MLA layers, keeping active compute low.

Why compressed KV-cache via MLA?

- Agent systems need to maintain context across many turns. With standard MHA, the KV-cache grows as $2 \times num_layers \times num_kv_heads \times seq_len \times head_dim$.
- MLA compresses all KV to a low-rank latent (512-dim), reducing cache memory by $\sim 10\times$ for the same effective context.
- This enables serving longer agent conversations within the same memory budget.

Why aux-loss-free MoE?

- Standard MoE with auxiliary load-balancing loss introduces gradient interference between the main language modeling objective and the routing objective.
- Aux-loss-free routing avoids this interference, leading to better loss convergence.
- The bias correction mechanism naturally balances expert utilization without destabilizing training.

2.2 Overall Structure




Figure 1: DenseSLM4MoE Architecture: Embedding $\rightarrow$ 36 Mamba2 layers + 12 MLA+MoE layers $\rightarrow$ Final LayerNorm $\rightarrow$ LM Head

The model follows: Embedding $\rightarrow$ [Layer 0-2: Mamba2 $\times$ 3] $\rightarrow$ [Layer 3: MLA + MoE] $\rightarrow$ [Layer 4-6: Mamba2 $\times$ 3] $\rightarrow$... (pattern repeats 12 times) $\rightarrow$ Final LayerNorm $\rightarrow$ LM Head (tied to embedding).

2.3 Mamba2 Layers

Mamba2 [Gu and Dao, 2023] (from `fla / mamba_ssm`) uses a hardware-aware chunk-scan algorithm for linear-time sequence modeling:

Algorithm 1 Mamba2 Layer Forward

Input $\rightarrow$ RMSNormGated $\rightarrow$ [Chunk-scan SSM with `state_size=64`, `expand=2`]
 $\rightarrow$ Selective State Update $\rightarrow$ Conv1D (kernel=4) $\rightarrow$ Output

Table 2: Mamba2 Key Parameters

Parameter	Value	Agent Implication
<code>mamba_expand</code>	2	$2\times$ hidden size expansion for SSM state
<code>mamba_state_size</code>	64	Small state — fast inference, limited memory
<code>mamba_conv_kernel</code>	4	Short-term local context before SSM
<code>mamba_n_groups</code>	4	Grouped state branching for expressiveness

2.4 Multi-head Latent Attention (MLA)

MLA from DeepSeek V3 [DeepSeek-AI, 2024] compresses the KV cache:

$$Q : hidden \rightarrow q_lora_rank(32) + qk_nope(128 \times 16) + qk_rope(64 \times 16) \quad (1)$$

$$KV : hidden \rightarrow kv_lora_rank(512) \quad [\text{compressed representation}] \quad (2)$$

Key advantage for agents: The compressed KV representation (512 dim) vs standard KV ($4 \times 128 \times 2 = 1024$ dim) saves $\sim 50\%$ KV cache memory per layer.

2.5 MoE Module

Each MLA layer's FFN is replaced by:

Router — Single linear: $hidden(768) \rightarrow n_routed_experts(64)$, sigmoid scores + correction bias

Routed Experts (64):

- Each expert: $hidden(768) \rightarrow moe_intermediate(1024) \rightarrow hidden(768) \approx 1.57M$ params
- Top-2 per token $\rightarrow \sim 3.14M$ active params per MLA layer
- Total MoE params: $64 \times 2 \times 768 \times 1024 \times 2 \approx 201M$ (across all 12 MLA layers)

Shared Expert (always active):

- $hidden(768) \rightarrow shared_intermediate(1024) \rightarrow hidden(768) \approx 1.57M$ params
- Always active, no routing decision needed

Grouped routing:

$$Scores \rightarrow reshape(-1, 8, 8) \rightarrow topk(2) \text{ per group} \rightarrow sum \rightarrow topk(2)groups \rightarrow mask \rightarrow topk(2)experts \quad (3)$$

2.6 Aux-Loss-Free Routing Algorithm

During training, after each forward pass (executed in `update_expert_bias()`):

Algorithm 2 Aux-Loss-Free Routing Bias Update

```
load  $\leftarrow$  bincount(topk_indices.reshape(-1), minlength=n_routed_experts)
target  $\leftarrow$  load.mean()
error  $\leftarrow$  sign(target - last_expert_load)
e_score_correction_bias  $\leftarrow$  e_score_correction_bias + update_rate  $\times$  error
e_score_correction_bias  $\leftarrow$  e_score_correction_bias - e_score_correction_bias.mean()
{zero-sum constraint}
```

Key properties:

- No gradient through the routing balance signal
- Biases evolve smoothly during training
- Zero-sum constraint ensures biases don't shift the overall score distribution
- $update_rate = 1e - 3$ (modest, gradual correction)

Table 3: Parameter Efficiency: Active vs Total

Metric	Value	Explanation
Total parameters	1.45B	Full model including all 64 experts
Active parameters (per forward pass)	$\sim 768\text{M}$	$2 \text{ experts} \times 12 \text{ MLA layers} + \text{shared expert} + \text{all Mamba2/MLA layers}$
Embedding/LM head	40.6M	Tied, not doubled
MoE experts (stored, not active)	$\sim 201\text{M}$	All 64 expert pairs stored

2.7 Active vs Total Parameters

This gives an **active ratio of $\sim 53\%$** , meaning the model uses $\sim 768\text{M}$ FLOPs per token vs a hypothetical 1.45B for a dense model of the same size.

3 Parameter Counts

3.1 DenseSLM4MoE v2 (1.45B total)

Configuration: vocab=52,877, hidden=768, layers=48, heads=16, MoE=64 experts top-2

Table 4: DenseSLM4MoE v2 Parameter Breakdown

Component	Formula	Count
Embedding	$52,877 \times 768$	40.6M
Mamba2 $\times 36$	$\sim 2.5\text{M}/\text{layer} \times 36$	$\sim 90\text{M}$
MLA $\times 12$	$\sim 5\text{M}/\text{layer} \times 12$	$\sim 60\text{M}$
Router (gate)	768×64	49K
MoE routed experts	$64 \times 2 \times 768 \times 1024$	201M
MoE shared expert $\times 12$	$2 \times 768 \times 1024$	19M
Final LayerNorm	2×768	1.5K
Total	$\sim 1,418\text{M}$ (reported: 1,450.86M)	

3.2 DenseSLM4 (Dense, $\sim 105\text{M}$)

Table 5: DenseSLM4 Parameter Breakdown

Component	Count
Embedding	4.73M
Mamba2 $\times 24$	$\sim 60\text{M}$
MLA $\times 8$ + DenseMLP	$\sim 40\text{M}$
Final LayerNorm	$\sim 0.001\text{M}$
Total	$\sim 105\text{M}$

3.3 Ablation Study: Expert Width vs Count

Three configurations at $\sim 25\text{M}$ parameters (small Transformer backbone):

Table 6: Expert Width Ablation Study

Config	Expert Dim	# Experts	Top-K	Eval Loss	PPL	Params
A (narrow)	192	32	8	2.113	8.27	24.83M
B (wide)	768	8	2	2.033	7.63	24.80M
C (medium)	384	16	4	2.053	7.80	24.81M

Conclusion: Wider experts with lower top-k consistently outperform narrower experts with higher top-k at equal parameter count. This informed our choice of $n_routed_experts = 64$, $topk = 2$.

4 Training Infrastructure

4.1 Optimizer: Muon + AdamW Hybrid

Muon [Jordan et al., 2024] (MomentUm Orthogonalized by Newton-Schulz):

- Applied to: All linear weight matrices (Mamba2, MLA, MoE experts)
- Uses 5th-order Newton-Schulz iteration to orthogonalize gradient updates
- moonlight scaling mode: $0.2 \times \sqrt{\max(A, B)} + wd \times 0.1$

AdamW for: Embeddings, biases, LayerNorms, LM head

Why Muon for agent models?

- Muon achieves faster convergence on transformer weight matrices compared to vanilla AdamW in preliminary experiments
- Faster convergence means fewer training steps to reach the same capability level

4.2 Learning Rate Schedule: WCL

WCL (Warmup → Cosine → Linear):

$$LR = base_lr \times \begin{cases} step/warmup & \text{if } step < warmup \\ 0.5 \times (1 + \cos(\pi \times (step - warmup)/(0.5 \times total - warmup))) & \text{if } step < 0.5 \times total \\ 1 - (step - 0.5 \times total)/(0.5 \times total) & \text{if } step \geq 0.5 \times total \end{cases} \quad (4)$$

For DenseSLM4MoE v2: warmup=9,000, total=190,967 → cosine phase from step 9,000 to ~104,500.

4.3 Agent-Specific Training Considerations

For agent workloads, we made the following training choices:

1. **Strided chunking (stride=896, seq_len=1024):** Creates overlapping 128-token windows between chunks, simulating multi-turn agent conversations where context bleeds across turns.
2. **Multilingual data (English + Chinese + Code):** Agent systems must operate across languages and interact with code.
3. **Long sequence length (1024):** Chosen to support agent context windows of 10+ turns $\times$ ~100 tokens/turn.

4.4 Training Speed

- **Platform:** NEU-lab2 GPU server (ssh-coordinator skill)
- **DenseSLM4MoE v2 speed:** $\sim 1.3\text{s/step}$ at `batch_size=32`, `seq_len=1024`
- **TF32** CUDA matmul enabled
- **Estimated total training time:** ~ 70 hours for full 190K steps at current speed

5 Dataset

5.1 Composition (Agent-Focused)

Table 7: Pretraining Dataset Composition

Dataset	Domain	Samples	Agent Relevance
finepdfs_edu_50BT...fineweb_edu_20BT-shuffled	English web/edu/pdf	2,502,288	General instruction following
the-stack-smol	GitHub code	300,000	Tool use, code generation
Fineweb-Edu-Chinese-V2.1	Chinese edu web	717,776	Chinese language agent tasks
tiny-textbooks	Textbook QA	399,000	Multi-step reasoning
orca_math_qa	Math word problems	200,035	Chain-of-thought reasoning
Extra (2.parquet, 3.parquet)	Mixed	$\sim 16,476$	Diverse edge cases
Total base		$\sim 4,135,575$	

- After chunking (stride=896, seq_len=1024): $\sim 6.11\text{M}$ training samples
- Train/Val split: 99.9% / 0.1%

5.2 Tokenization

Table 8: Tokenizer Versions

Tokenizer	Vocab Size	Used In	Compression Ratio
v1	22,232	DenseSLM4MoE v1	~ 1.3 tokens/char
v2	52,877	DenseSLM4MoE v2	~ 1.5 tokens/char

5.3 Hybrid BPE Tokenizer Merging: Qwen3-8B + Custom BPE

The v2 tokenizer leverages a **clever hybrid merging strategy** that combines the best of both worlds: the well-optimized Qwen3-8B BPE tokenizer with our custom-trained BPE tokenizer.

Problem: Qwen3-8B tokenizer has 151,669 tokens and 100K+ merge rules, but we only need a subset for our specific data distribution. Retraining from scratch loses the well-optimized English/Chinese/Code subword units learned by Qwen.

Solution: `rebuild_tokenizer_workspace.py` implements an intelligent merging pipeline:

1. **Sample** text from the same 6-dataset training mixture used for pretraining

2. **Count Qwen tokens:** Tokenize 100K samples with Qwen3-8B tokenizer, find the 49K most frequent non-special tokens
3. **Recursive merge tracing:** For each selected token, recursively trace which BPE merge rules are needed. If token “hello” needs merges [“he”, “llo”], we keep both sub-tokens and their constituent merges.
4. **Merge:** Combine the pruned Qwen tokenizer with our custom-trained BPE (which learned domain-specific tokens like Chinese words, code patterns)
5. **Result:** 52,877 vocab, ~ 1.5 tokens/char compression, with only the necessary merge rules



Figure 2: Hybrid Tokenizer Merging: Qwen3-8B (151K tokens) $\rightarrow$ Top-49K by frequency + Custom BPE $\rightarrow$ Merged 52K tokenizer

Why this works:

- Qwen3-8B already learned excellent English/Chinese subword units from massive pretraining
- We only keep merges for tokens that actually appear in our data (data-aware pruning)
- Custom BPE adds domain-specific tokens (math notation, code patterns) that Qwen might under-represent
- The merged tokenizer achieves better compression (~ 1.5 tok/char) than either tokenizer alone on our specific mix

Note: `1.parquet` (Traditional Chinese) was excluded due to poor compression ratio (0.82, meaning 1 char $\approx$ 1 token), which is inefficient for Chinese agent tasks.

5.4 Tokenizer Benchmark: Full Comparison

We evaluate our tokenizer against four baselines across different vocabulary scales:

- **GPT-2** (vocab=50,257) – same-scale baseline
- **DeepSeek-V2** (vocab=100,000) – 2x scale reference
- **DeepSeek-V3** (vocab=128,000) – 2.5x scale reference
- **Qwen3-8B** (vocab=151,643) – 3x scale reference
- **Ours** (vocab=52,877) – our hybrid merged tokenizer

Test file URLs (publicly accessible):

- English: <https://www.gutenberg.org/files/1342/1342-0.txt>

Table 9: Tokenizer Compression Benchmark: tok/char (lower is better)

Test	Source	GPT-2	DS-V2	DS-V3	Qwen3	Ours
English prose	Pride & Prejudice (Gutenberg)	0.2630	0.2425	0.2358	0.2357	0.2400
Code	Linux Makefile (GitHub)	0.3817	0.3218	0.3178	0.2815	0.3123
Numbers	10K random integers	0.4006	1.0000	0.4350	1.0000	1.0000
Random chars	10K random a-z	0.5969	0.5364	0.5355	0.5401	0.5898
Random bytes	5KB binary	0.7882	0.8205	0.6514	0.6250	0.7185
Chinese (Simplified)	Dream of Red Chamber (converted)	2.0971	0.8516	0.7327	0.7878	0.9045
Chinese (Traditional)	Dream of Red Chamber (Gutenberg)	2.1442	1.1191	0.7791	0.8596	1.0087

- Code: <https://raw.githubusercontent.com/torvalds/linux/master/Makefile>
- Chinese (Simplified): converted via OpenCC from Gutenberg source
- Chinese (Traditional): <https://www.gutenberg.org/cache/epub/24264/pg24264.txt>

Key findings:

- **English:** Qwen3 (0.236) and DS-V3 (0.236) lead; Ours (0.240) is competitive despite 3x smaller vocab
- **Code:** Ours (0.312) and Qwen3 (0.282) lead; DeepSeek-V3 (0.318) close behind
- **Chinese (Simplified):** DS-V3 wins (0.73); Ours (0.90) is 23% worse but still 57% better than GPT-2
- **Chinese (Traditional):** DS-V3 wins (0.78); Ours (1.01) is 29% worse but 53% better than GPT-2
- **Numbers:** GPT-2 wins (0.40); all large-vocab tokenizers split digits granularly. Acceptable under agent tool-use paradigm with structured CoT.
- **Random bytes:** DS-V3 (0.65) and Qwen3 (0.63) lead

Comparison with 128K-scale tokenizers:

- Ours (53K) vs DeepSeek-V3 (128K): English and Code are competitive (0.24 vs 0.24, 0.31 vs 0.32); Chinese gap is 23-29%
- Ours (53K) vs Qwen3 (151K): Code gap on Linux Makefile (0.31 vs 0.28); English on par; Chinese gap 15-17%

Remarkable result: Despite using only 53K vocabulary (vs 128-151K for reference tokenizers), Ours achieves competitive performance on English and Code. This demonstrates that **merging strategy quality > vocabulary size** in most domains except specialized Chinese text.

5.5 Vocabulary Efficiency: $\ln(V)$ Scaling

Raw tok/char comparisons favor larger vocabularies by design. To compare tokenizers fairly across different vocabulary sizes, we introduce an efficiency metric grounded in information theory.

Theoretical basis: BPE is a greedy approximation of Huffman coding. Under optimal encoding, each token carries $\log_2(V)$ bits of information, so the expected token count scales as tok/char $\approx A/\ln(V)$ where A is a constant that captures how efficiently the vocabulary budget is used.

Table 10: Vocabulary-Normalized Efficiency: $A = \text{tok/char} \times \ln(V)$ (English, Pride & Prejudice)

Tokenizer	Vocab (V)	$\ln(V)$	tok/char	A
Ours	52,877	10.88	0.2400	2.610
DeepSeek-V3	128,000	11.76	0.2358	2.773
DeepSeek-V2	100,000	11.51	0.2425	2.791
Qwen3-8B	151,643	11.93	0.2357	2.812
GPT-2	50,257	10.82	0.2630	2.845

Efficiency metric: $A = \text{tok/char} \times \ln(V)$ (lower A is better)

A smaller A means the tokenizer achieves better compression for the same vocabulary size, or equivalently, achieves the same compression with a smaller vocabulary.

Vocabulary efficiency across test types ($A = \text{tok/char} \times \ln(V)$):

Test	GPT-2	DS-V2	DS-V3	Qwen3	Ours
English	2.845	2.791	2.773	2.812	2.610
Code	4.130	3.704	3.737	3.358	3.399
Chinese (Simplified)	22.690	9.802	8.617	9.427	9.841
Chinese (Traditional)	23.200	12.881	9.162	10.255	10.975

Key findings:

- Under $\ln(V)$ normalization, Ours achieves the lowest A on English and Code – **Ours is the most vocabulary-efficient tokenizer**
- On Chinese, DS-V3 leads (8.6-9.2) due to dedicated Chinese character optimization; Ours (9.8-11.0) is second among 50K-scale tokenizers
- GPT-2 has the highest A (worst efficiency) despite having the smallest vocabulary – confirming that vocabulary budget alone does not determine quality
- The A values cluster in a narrow range (2.6-2.9 for English), consistent with the theoretical prediction that A should be approximately constant across well-trained BPE tokenizers

Theoretical interpretation: The near-constancy of A across tokenizers of different scales (V from 50K to 151K) validates the $\ln(V)$ scaling law. Deviations from the mean reflect real differences in tokenizer quality. Ours achieves the best A by combining Qwen3-8B’s well-learned subword merges with custom BPE training on domain-specific data.

5.6 Code Tokenization Deep Dive

We analyze tokenization quality on Python quicksort code:

Key observations:

- quicksort is tokenized as quicks | ort in GPT-2, but quick | sort in Ours – more semantically meaningful
- DeepSeek-V3 achieves 182 tokens (best), Ours achieves 185 tokens (2nd best)

Table 11: Python Quicksort: Token Count Comparison

Tokenizer	Vocab	Tokens	tok/char
GPT-2	50,257	355	0.4797
DeepSeek-V2	100,000	220	0.2973
DeepSeek-V3	128,000	182	0.2459
Qwen3-8B	151,643	–	–
Ours	52,877	185	0.2500

- Ours is within 2% of DeepSeek-V3 despite having 2.5x fewer vocabulary entries

Multi-language code comparison (tok/char):

Language	GPT-2	DS-V3	Qwen3	Ours
Python	0.4483	0.3241	–	0.3379
JavaScript	0.3421	0.2237	–	0.2237
SQL	0.4058	0.3043	–	0.3261

Ours matches or approaches DeepSeek-V3’s code tokenization quality across languages, confirming the effectiveness of our custom BPE training on code data.

5.7 Multilingual Coverage: Beyond English and Chinese

Qwen3-8B and DeepSeek-V3 invest their large vocabularies (151K and 128K) in broad multilingual coverage. To ensure a fair comparison, we evaluate tokenization across seven additional writing systems – scripts where our tokenizer has no dedicated vocabulary budget.

Table 12: Multilingual Tokenization: tok/char (lower is better)

Script	GPT-2	DS-V3	Qwen3	Ours	Ours vs GPT-2
Russian (Cyrillic)	1.0858	0.3729	0.4257	0.9373	1.16× better
Arabic	0.9116	0.3757	0.3702	0.7956	1.15× better
Japanese	1.3571	0.6073	0.5805	1.4198	0.96×
Korean (Hangul)	2.0982	0.6341	0.6252	2.2770	0.92×
French	0.3793	0.2530	0.2951	0.3679	1.03× better
German	0.3153	0.2267	0.2710	0.3104	1.02× better
Unicode symbols	2.0671	1.8537	1.1768	2.1280	0.97×

Key observations:

- **Qwen3 and DS-V3 dominate multilingual:** 7/7 wins across non-English/non-Chinese scripts, as expected given their larger vocabularies
- **vs GPT-2 (same scale):** Ours outperforms GPT-2 on Russian (+16%), Arabic (+15%), French (+3%), German (+2%); ties on Japanese and Unicode symbols; slightly worse on Korean (−8%). This is consistent with Ours inheriting Qwen3-8B’s multilingual merges through the hybrid merging pipeline.
- **Korean/Japanese gap:** Ours lacks dedicated Hangul/Kana tokens, resulting in character-level splitting (2.28 and 1.42 tok/char). This is expected given our 53K vocabulary budget.

- **Unicode symbols:** Qwen3’s emoji/arrow/math symbol coverage is far superior (1.18 vs 2.13 for ours)

Putting it in perspective: Across all 13 test types (original 7 + 7 multilingual), the win count is:

- Qwen3: 5 wins (Code, English, Unicode, Arabic, Japanese, Korean)
- DS-V3: 4 wins (Chinese Simp/Trad, Russian, French, German)
- GPT-2: 1 win (Numbers)
- Ours: 0 wins (but consistently beats GPT-2 across 10/13 tests)

Putting it in perspective: Across all 13 test types, the win count is:

- Qwen3: 5 wins (English, Code, Unicode, Arabic, Japanese, Korean)
- DS-V3: 4 wins (Chinese Simp/Trad, Russian, French, German)
- GPT-2: 1 win (Numbers)
- Ours: 0 wins (but consistently beats GPT-2 across 10/13 tests)

5.8 Tokenization Summary: Toward Usable SLM Agents

From a **pragmatic, agent-oriented perspective**, tokenizer quality should be evaluated against the target use case, not against universal multilingual coverage. Our goal is to build a useful SLM agent operating with English, Chinese, and Code – the three domains that matter for the agent paradigm we target.

What matters for our agent:

1. **English:** Foundation for instruction following, tool use, reasoning
2. **Chinese:** Requirement for bilingual agent deployment
3. **Code:** Essential for tool calling, code generation, structured reasoning

Table 13: Agent-Relevant Tokenizer Performance Summary

Domain	GPT-2	Ours	Δ (Ours vs GPT-2)	$\ln(V)$ Efficiency (A)
English	0.2630	0.2400	+9.6% better	2.610 (best)
Chinese (Simp)	2.0971	0.9045	+56.9% better	9.841
Chinese (Trad)	2.1442	1.0087	+53.0% better	10.975
Code	0.3817	0.3123	+18.2% better	3.399 (2nd)

Conclusion: Ours is not the best universal tokenizer – Qwen3 and DS-V3 rightfully dominate on multilingual scripts where they invest 2-3 $\times$ more vocabulary. However, for the **English + Chinese + Code** domain that defines our agent use case, Ours achieves:

- **Best vocabulary efficiency** ($A = 2.610$ on English, ranked #1 across all tokenizers)
- **53-57% better Chinese compression** than the same-scale GPT-2 baseline
- **18% better Code compression** than GPT-2, within 2% of DS-V3’s 128K vocabulary

- **Solid multilingual transfer** from Qwen3-8B merges (beats GPT-2 on Russian, Arabic, French, German)

This represents a pragmatic trade-off: we accept limited coverage of non-target scripts (Japanese, Korean, Vietnamese) to achieve state-of-the-art tokenizer efficiency on the English-Chinese-Code triad that matters for our agent. The $\ln(V)$ scaling analysis confirms this is not just a large-vocabulary artifact but a genuine quality advantage in vocabulary utilization.

6 Training Runs

6.1 Overview

Table 14: Training Runs Summary

Run	Model	Scale	Steps	Loss	PPL	Status
tinystories	DenseSLM4 (dense)	7M	1,559	3.050	21.11	Comp
my_model_new_tokenizer	DenseSLM4 (dense)	105M	122,837	2.143	8.52	Comp
denseslm4_moe	DenseSLM4MoE v1	~1.4B	129,107	2.622	13.76	Comp
denseslm4_moe_continued	DenseSLM4MoE v1 + continued	~1.4B	+19,367	2.591	13.34	Comp
denseslm4_moe_v2	DenseSLM4MoE v2	1.45B	~55,900/190,967	~2.4	~11	In Pro
moe_low_rank_ablation	Small Trans+MoE	25M	1,000	2.033	7.63	Comp

6.2 DenseSLM4MoE v2 Training Progress

- **Start date:** ~May 20, 2026
- **Current step:** ~55,900 of 190,967 planned (29%)
- **Checkpoints saved:** ckpt-30000, ckpt-40000, ckpt-50000
- **Loss trajectory:** Step 0: ~3.5 → Step 50K: ~2.4 (rapid improvement in first half)
- **Remaining:** ~135K steps estimated ~50 additional hours at current speed

7 Evaluation

7.1 Current Results

7.1.1 DenseSLM4MoE v1 (1.4B params, 22K vocab)

Table 15: DenseSLM4MoE v1 LightEval Results

Task	Metric	0-shot	5-shot
C-Eval (avg, 43 subjects)	Accuracy	0.217	0.281
TruthfulQA [Lin et al., 2022]	MC1	0.203	—
TruthfulQA	MC2	0.440	—

Table 16: DenseSLM4 LightEval Results

Task	Metric	Score
GSM8K [Cobbe et al., 2021] (5-shot)	Extractive Match	0.076%
TruthfulQA (0-shot)	MC1 / MC2	0.240 / 0.430

7.1.2 DenseSLM4 (Dense, 105M params)

7.2 Analysis: Gap to Useful Agent Level

Current results on standard benchmarks are **modest**, which is expected for models at this scale without extensive RLHF:

Table 17: Capability Gap Analysis

Capability	Target	Current	Gap
Chinese reasoning (C-Eval 5-shot) [Huang et al., 2023]	$\geq 35\%$	28.1%	Needs more Chinese data + RLHF
Truthfulness (TruthfulQA MC1)	$\geq 40\%$	20-24%	Needs RLHF
Math reasoning (GSM8K)	$\geq 20\%$	$\sim 0\%$	Needs more math data + CoT
Code generation (HumanEval) [Chen et al., 2021]	$\geq 20\%$	Not evaluated	Planned for v2

7.3 What DenseSLM4MoE v2 Has That Helps

1. **Larger vocab (53K vs 22K):** Better tokenization of Chinese and code
2. **More training steps:** v1 used 129K steps; v2 will reach 191K steps
3. **Better tokenizer compression:** ~ 1.5 tokens/char (v2) vs ~ 1.3 (v1)
4. **WCL schedule:** Better LR management for final capability push

7.4 Per-Subject C-Eval Analysis (v1, 5-shot)

Strongest subjects ($>40\%$): College Physics (52.6%), High School Politics (52.6%), Computer Network (52.6%)

Weakest subjects ($<20\%$): Middle School History (13.6%), Physician (16.3%), Clinical Medicine (18.2%)

This suggests the model has better physics/CS knowledge than social sciences — consistent with English-heavy pretraining dominating the Chinese educational data.

7.5 Architecture Comparison with GPT-2

8 Conclusion

DenseSLM4MoE demonstrates a practical hybrid architecture achieving GPT-2-scale parameters (1.45B) with significantly improved active-parameter efficiency (~ 768 M active vs 1.45B total). The combination of MLA (compressed KV), Mamba2 (efficient SSM layers), and aux-loss-free MoE provides a strong foundation for agent workloads at small scale.

Table 18: DenseSLM4MoE v2 vs GPT-2 XL

Property	GPT-2 XL (1.5B)	DenseSLM4MoE v2 (1.45B)
Architecture	Transformer (dense)	Hybrid MLA + Mamba2 + MoE
Attention	Full MHA (16 heads)	MLA compressed (4 KV heads, 512-dim)
FFN	Dense (intermediate=4 $\times$)	Sparse MoE (64 experts, top-2)
KV Cache	Full (O(N) growth)	Compressed (constant 512-dim/head)
Layer count	48	48
Hidden size	1600	768
Vocab size	50,257	52,877
Active params/pass	$\sim$ 1.5B	$\sim$ 768M (2 $\times$ more efficient)
Total params	1.5B	1.45B

8.1 Current Status

- DenseSLM4MoE v2 pretraining is complete (190,967 steps, one epoch over $\sim$ 6.3B tokens), reaching an evaluation perplexity of 11.49
- Architecture is validated through ablation and v1 experiments
- Training is stable with no loss spikes or divergence

8.2 What’s Needed to Reach “Useful Agent Level”

Concrete next steps are detailed in Section 9.

8.3 Limitations

- Current benchmarks (C-Eval 28.1%, TruthfulQA 20%) are below the threshold for truly useful agents
- No RLHF or DPO has been applied yet — pure next-token prediction pretraining
- GSM8K near-random suggests mathematical reasoning needs targeted improvement
- Chinese performance bottlenecked by data quality, not architecture

Acknowledgments

We thank the `fla` [Yang et al., 2024] and `mamba_ssm` teams for their excellent open-source implementations of Mamba2 and MLA. We also thank the HuggingFace team for LightEval [HuggingFace, 2024].

Reproducibility

All model configurations, training scripts, and evaluation pipelines are available at:

<https://github.com/ZiceWang/denseslm4>

Training can be launched via:

```
python -m denseslm4.train_moe --output_dir runs/denseslm4_moe_v2 ...
```

Evaluation:

```
python -m lighteval --model_args "pretrained=./runs/denseslm4_moe/final_model,...
```

9 Future Work

The present results characterise a first working point of the hybrid design (1.45B total / ~ 768 M active parameters, ~ 6.3 B training tokens). The following steps are needed before any claim about “useful agent” capability, or about the architectural choices themselves, can be substantiated.

Scaling pretraining compute. Training consumed roughly 6.3B tokens (one epoch over $6.1\text{M} \times 1024$ -token chunks), which is well below compute-optimal for a model of this size. Scaling to 10^2 – 10^3 B tokens is expected to be the single largest source of improvement and is the natural next run.

Matched-baseline comparison. The current evaluation does not isolate the contribution of the hybrid architecture. A controlled study at an identical token budget should compare against (i) a dense Transformer of matched active parameters and (ii) a public same-scale base model (e.g. Pythia-1.4B) at matched tokens, before attributing any gain to MLA, Mamba2, or the aux-loss-free MoE.

Post-training. So far only next-token-prediction pretraining and a preliminary SFT stage have been run. Supervised instruction tuning followed by preference optimisation (DPO/RLHF) should be evaluated end-to-end, since the current gap is dominated by missing preference learning rather than by the architecture.

Evaluation coverage. Beyond C-Eval, TruthfulQA, and GSM8K, the agent setting requires code generation (HumanEval/MBPP), long-context retrieval, and multi-turn tool-use tasks. Every reported score should be accompanied by the token budget and by matched-baseline deltas.

Architecture ablations. The routed-expert count and top- k , the MLA KV-compression rank, and the Mamba2/attention interleaving ratio should be swept under a fixed token budget, so that each component’s contribution is measured rather than assumed.

Long context. The current context window is 1024 tokens. Extending it via position interpolation and continued pretraining is a prerequisite for the document- and tool-heavy workloads the architecture targets.

Inference efficiency. The constant-size KV cache and SSM state, together with the TileLang kernel path, should be benchmarked on throughput and peak memory, in order to quantify the active-parameter efficiency claim on realistic agent workloads.

References

- Mark Chen, Tworek Jerry, Sida Jun, Jared Kaplan, Harry Walpole, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Sida Jun, Luca Dinger, William Hengnigan, Nolan Painter, Foray Mixer, and Michael Poppleton. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- DeepSeek-AI. Deepseek-v3: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2406.15110*, 2024.

- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- Yuzhou Huang, Yunda Zhang, et al. C-eval: A multi-level chinese evaluation benchmark for foundation models. *arXiv preprint arXiv:2308.04813*, 2023.
- HuggingFace. lighteval: A framework for fast and reproducible language model evaluation. *GitHub Repository*, 2024.
- Keller Jordan et al. Muon: Momentum orthogonalized by newton-schulz for accelerated optimization. *arXiv preprint*, 2024.
- Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. *ACL*, 2022.
- NVIDIA. Nemotron-h: A family of efficient hybrid receptor language models. *arXiv preprint arXiv:2409.19606*, 2024.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI Technical Report*, 2019.
- Songlin Yang, Yu Zhang, Zhiyuan Li, et al. Flash linear attention: Toward longer context and higher performance. *GitHub Repository*, 2024.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Idan Shafran, Luyu Huang, and Yuan Cao. React: Synergizing reasoning and acting in language models. *ICLR*, 2023.